

Exercises 2

Flights at ABIA

Here we will be looking at flights into and out of Austin in the year 2008. I will present some interesting findings about flight delays out of Austin.

First, read in the data and examine some of the features.

```
atxflights = read.csv('ABIA.csv')
head(atxflights)
```

```
##   Year Month DayOfMonth DayOfWeek DepTime CRSDepTime ArrTime CRSArrTime
## 1 2008     1           1           2     120      1935      309      2130
## 2 2008     1           1           2     555        600      826      835
## 3 2008     1           1           2     600        600      728      729
## 4 2008     1           1           2     601        605      727      750
## 5 2008     1           1           2     601        600      654      700
## 6 2008     1           1           2     636        645      934      932
##   UniqueCarrier FlightNum TailNum ActualElapsedTime CRSElapsedTime AirTime
## 1              9E      5746  84129E              109             115      88
## 2              AA      1614  N438AA              151             155     133
## 3              YV      2883  N922FJ              148             149     125
## 4              9E      5743  89189E              86             105      70
## 5              AA      1157  N4XAAA              53              60      38
## 6              NW      1674  N967N              178             167     145
##   ArrDelay DepDelay Origin Dest Distance TaxiIn TaxiOut Cancelled
## 1       339       345   MEM  AUS     559      3      18         0
## 2        -9        -5   AUS  ORD     978      7      11         0
## 3         -1         0   AUS  PHX     872      7      16         0
## 4       -23        -4   AUS  MEM     559      4      12         0
## 5         -6         1   AUS  DFW     190      5      10         0
## 6          2        -9   AUS  MSP    1042     11      22         0
##   CancellationCode Diverted CarrierDelay WeatherDelay NASDelay
## 1                  0          339              0          0
## 2                  0           NA              NA         NA
## 3                  0           NA              NA         NA
## 4                  0           NA              NA         NA
## 5                  0           NA              NA         NA
## 6                  0           NA              NA         NA
##   SecurityDelay LateAircraftDelay
## 1              0                 0
## 2             NA                 NA
## 3             NA                 NA
## 4             NA                 NA
## 5             NA                 NA
## 6             NA                 NA
```

```
summary(atxflights)
```

```
##      Year      Month      DayOfMonth      DayOfWeek
## Min.   :2008   Min.   : 1.00   Min.   : 1.00   Min.   :1.000
## 1st Qu.:2008   1st Qu.: 3.00   1st Qu.: 8.00   1st Qu.:2.000
```

```

## Median :2008      Median : 6.00      Median :16.00      Median :4.000
## Mean :2008      Mean : 6.29      Mean :15.73      Mean :3.902
## 3rd Qu.:2008      3rd Qu.: 9.00      3rd Qu.:23.00      3rd Qu.:6.000
## Max. :2008      Max. :12.00      Max. :31.00      Max. :7.000
##
##      DepTime      CRSDepTime      ArrTime      CRSArrTime
## Min. : 1      Min. : 55      Min. : 1      Min. : 5
## 1st Qu.: 917      1st Qu.: 915      1st Qu.:1107      1st Qu.:1115
## Median :1329      Median :1320      Median :1531      Median :1535
## Mean :1329      Mean :1320      Mean :1487      Mean :1505
## 3rd Qu.:1728      3rd Qu.:1720      3rd Qu.:1903      3rd Qu.:1902
## Max. :2400      Max. :2346      Max. :2400      Max. :2400
## NA's :1413      NA's :1567
## UniqueCarrier      FlightNum      TailNum      ActualElapsedTime
## WN :34876      Min. : 1      : 1104      Min. : 22.0
## AA :19995      1st Qu.: 640      N678CA : 195      1st Qu.: 57.0
## CO : 9230      Median :1465      N511SW : 180      Median :125.0
## YV : 4994      Mean :1917      N526SW : 176      Mean :120.2
## B6 : 4798      3rd Qu.:2653      N528SW : 172      3rd Qu.:164.0
## XE : 4618      Max. :9741      N520SW : 168      Max. :506.0
## (Other):20749      (Other):97265      NA's :1601
## CRSElapsedTime      AirTime      ArrDelay      DepDelay
## Min. : 17.0      Min. : 3.00      Min. : -129.000      Min. : -42.000
## 1st Qu.: 58.0      1st Qu.: 38.00      1st Qu.: -9.000      1st Qu.: -4.000
## Median :130.0      Median :105.00      Median : -2.000      Median : 0.000
## Mean :122.1      Mean : 99.81      Mean : 7.065      Mean : 9.171
## 3rd Qu.:165.0      3rd Qu.:142.00      3rd Qu.: 10.000      3rd Qu.: 8.000
## Max. :320.0      Max. :402.00      Max. : 948.000      Max. :875.000
## NA's :11      NA's :1601      NA's :1601      NA's :1413
##      Origin      Dest      Distance      TaxiIn
## AUS :49623      AUS :49637      Min. : 66      Min. : 0.000
## DAL : 5583      DAL : 5573      1st Qu.: 190      1st Qu.: 4.000
## DFW : 5508      DFW : 5506      Median : 775      Median : 5.000
## IAH : 3704      IAH : 3691      Mean : 705      Mean : 6.413
## PHX : 2786      PHX : 2783      3rd Qu.:1085      3rd Qu.: 7.000
## DEN : 2719      DEN : 2673      Max. :1770      Max. :143.000
## (Other):29337      (Other):29397      NA's :1567
##      TaxiOut      Cancelled      CancellationCode      Diverted
## Min. : 1.00      Min. :0.000000      :97840      Min. :0.000000
## 1st Qu.: 9.00      1st Qu.:0.000000      A: 719      1st Qu.:0.000000
## Median :12.00      Median :0.000000      B: 605      Median :0.000000
## Mean :13.96      Mean :0.01431      C: 96      Mean :0.001824
## 3rd Qu.:16.00      3rd Qu.:0.000000      3rd Qu.:0.000000
## Max. :305.00      Max. :1.000000      Max. :1.000000
## NA's :1419
##      CarrierDelay      WeatherDelay      NASDelay      SecurityDelay
## Min. : 0.00      Min. : 0.00      Min. : 0.00      Min. : 0.00
## 1st Qu.: 0.00      1st Qu.: 0.00      1st Qu.: 0.00      1st Qu.: 0.00
## Median : 0.00      Median : 0.00      Median : 2.00      Median : 0.00
## Mean :15.39      Mean : 2.24      Mean :12.47      Mean : 0.07
## 3rd Qu.:16.00      3rd Qu.: 0.00      3rd Qu.:16.00      3rd Qu.: 0.00
## Max. :875.00      Max. :412.00      Max. :367.00      Max. :199.00
## NA's :79513      NA's :79513      NA's :79513      NA's :79513
## LateAircraftDelay

```

```
## Min.   : 0.00
## 1st Qu.: 0.00
## Median : 6.00
## Mean   : 22.97
## 3rd Qu.: 30.00
## Max.   :458.00
## NA's   :79513
```

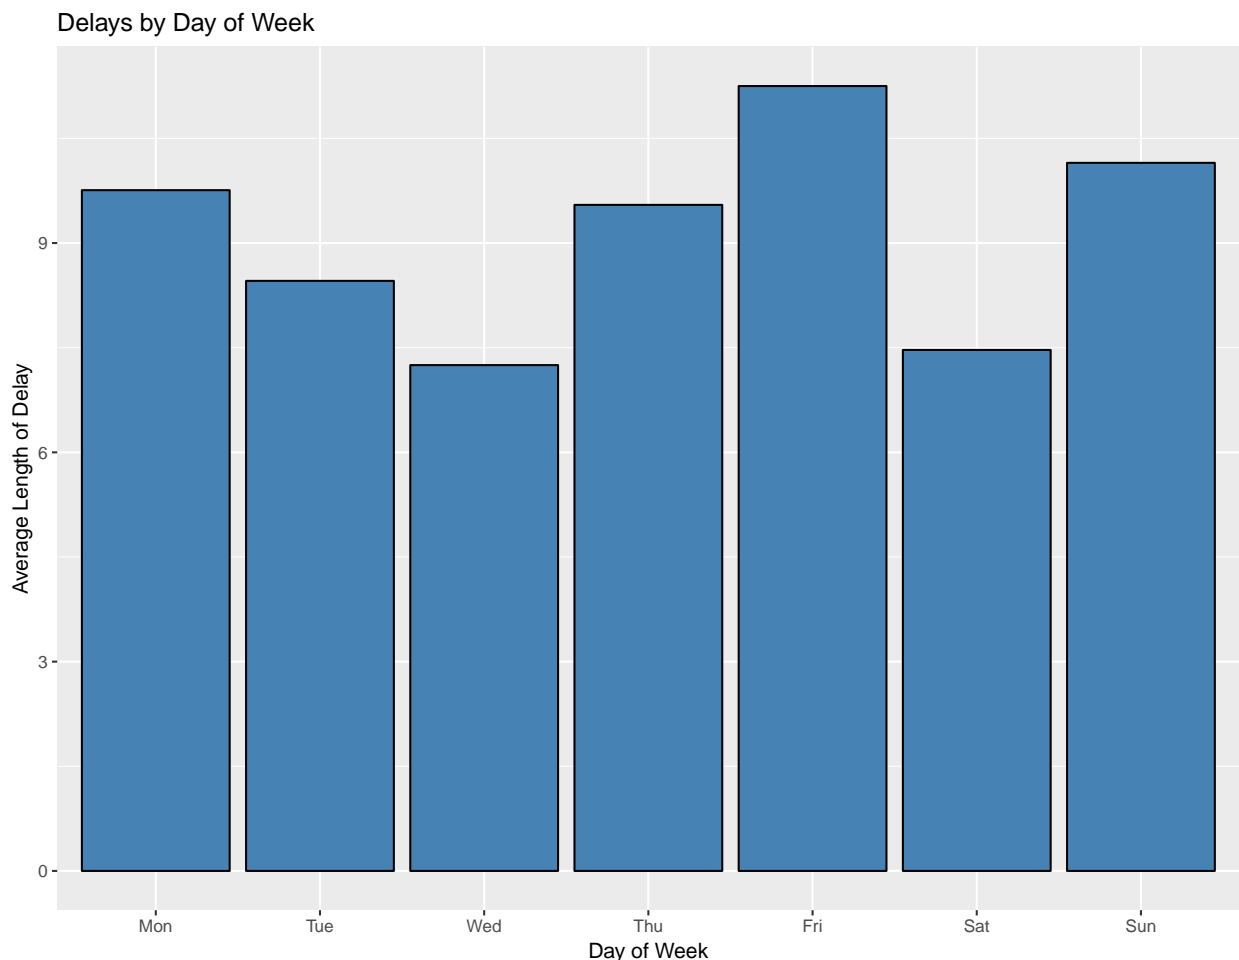
I want to take a look at flight delays by day of the week, and by month. However these variables are stored as numeric so we'll convert them to factors to make them better to deal with, as well as give them some labels.

```
# convert a couple variables to factors
atxflights$DayOfWeek = factor(atxflights$DayOfWeek, levels = c(1,2,3,4,5,6,7),
                              labels = c("Mon", "Tue", "Wed", "Thu", "Fri", "Sat", "Sun"))
atxflights$Month = factor(atxflights$Month, levels = c(1,2,3,4,5,6,7,8,9,10,11,12),
                          labels = c("Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec"))
```

I want to find the average flight delay length over each day of the week, so I'll store those in a new data frame.

```
# get the average length of delays by day of week
avgdelay = aggregate(atxflights$DepDelay, by = list(atxflights$DayOfWeek), FUN = mean, na.rm = TRUE)
avgdelay = as.data.frame(avgdelay)
```

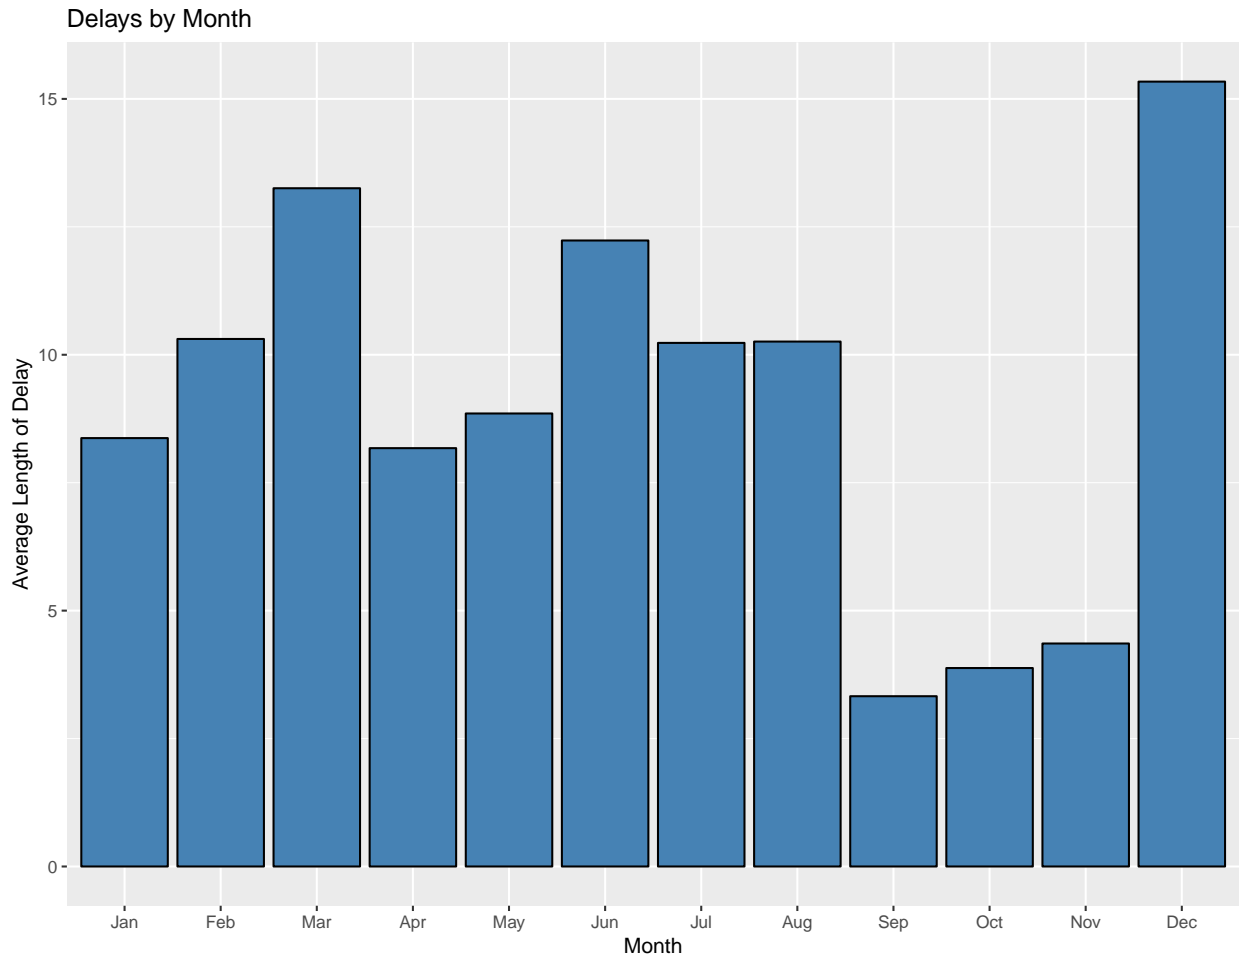
Now we plot flight delays by day of the week.



The worst day of the week for delays was Friday - this could be problematic as it is the end of the work week and many people would probably be traveling either back home or leaving on a trip. Wednesday and Saturday are pretty good days if you want to cut down on flight delay time.


Next I want to take a look at the worst delay times by month. Again we'll compute the averages and plot it.

```
# get average delay by month  
avgdelaymonth = aggregate(atxflights$DepDelay, by = list(atxflights$Month), FUN = mean, na.rm = TRUE)  
avgdelaymonth = as.data.frame(avgdelaymonth)
```



December is the worst month, followed by March - not surprisingly, as we have the holiday season in December as well as SXSW in Austin during the month of March - expect longer delays at those times.

Let's take a look at the month of December - when do you experience the worst delays?



December 27th had the worst delay times - just a couple days after Christmas. Looks like flying on Christmas Day is the way to go! Or do your traveling around December 8th and you'll experience almost no delays. Unfortunately school is still in session!

Author Attribution

Now we'll look at the C50 corpus and build a couple models to try to predict the author, and see how we do.

The first step was pre-processing which involved reading in the documents, cleaning up file names, getting author names from the files, building the corpus, and removing unnecessary things like numbers, punctuation, stop words, etc.

Then we build the Document Term Matrix and check the summary.

```
## <<DocumentTermMatrix (documents: 5000, terms: 45522)>>
## Non-/sparse entries: 954272/226655728
## Sparsity           : 100%
## Maximal term length: 45
## Weighting          : term frequency (tf)
```

There are 45000+ terms here - we need to remove sparse terms to cut that number down. We'll remove terms that didn't come up in 97.5% of the documents.

```
## <<DocumentTermMatrix (documents: 5000, terms: 1407)>>
```

```
## Non-/sparse entries: 581774/6453226
## Sparsity           : 92%
## Maximal term length: 18
## Weighting          : term frequency (tf)
```

Much better - about 1400 terms now.

Next we set up the training and test sets from the Document Term Matrix, and we will simply ignore words in the test set that didn't appear in the training set in order to avoid getting zero probabilities.

We will try a random forest using PCA to get the first 100 principal components to put in the random forest model.

The accuracy of the random forest turns out to be:

```
accuracy_rf
```

```
## [1] 0.51
```

Not so great as we get an accuracy of 51%. Let's try another model.

Next we'll try a Naive Bayes model. First we use the training set and apply a smoothing function, then get the predictions and check them against the test set.

Then we get the accuracy for Naive Bayes:

```
accuracy_nb
```

```
## [1] 0.6024
```

It turns out Naive Bayes has an accuracy of about 60%, better than that of random forest. The table below shows the authors with how well the model did at predicting each author.

##	author	nbcorrect	accuracy
## 1	AaronPressman	42	0.84
## 2	AlanCrosby	25	0.50
## 3	AlexanderSmith	21	0.42
## 4	BenjaminKangLim	11	0.22
## 5	BernardHickey	27	0.54
## 6	BradDorfman	44	0.88
## 7	DarrenSchuettler	14	0.28
## 8	DavidLawder	7	0.14
## 9	EdnaFernandes	20	0.40
## 10	EricAuchard	25	0.50
## 11	FumikoFujisaki	49	0.98
## 12	GrahamEarnshaw	41	0.82
## 13	HeatherScofield	17	0.34
## 14	JaneMacartney	23	0.46
## 15	JanLopatka	18	0.36
## 16	JimGilchrist	50	1.00
## 17	JoeOrtiz	34	0.68
## 18	JohnMastrini	40	0.80
## 19	JonathanBirt	32	0.64
## 20	JoWinterbottom	36	0.72
## 21	KarlPenhaul	47	0.94
## 22	KeithWeir	37	0.74
## 23	KevinDrawbaugh	27	0.54
## 24	KevinMorrison	27	0.54
## 25	KirstinRidley	34	0.68
## 26	KouroshKarimkhany	32	0.64

## 27	LydiaZajc	31	0.62
## 28	LynneO'Donnell	39	0.78
## 29	LynnleyBrowning	49	0.98
## 30	MarcelMichelson	30	0.60
## 31	MarkBendeich	20	0.40
## 32	MartinWolk	25	0.50
## 33	MatthewBunce	42	0.84
## 34	MichaelConnor	40	0.80
## 35	MureDickie	14	0.28
## 36	NickLouth	40	0.80
## 37	PatriciaCommins	31	0.62
## 38	PeterHumphrey	33	0.66
## 39	PierreTran	34	0.68
## 40	RobinSidel	40	0.80
## 41	RogerFillion	38	0.76
## 42	SamuelPerry	32	0.64
## 43	SarahDavison	27	0.54
## 44	ScottHillis	13	0.26
## 45	SimonCowell	27	0.54
## 46	TanEeLyn	21	0.42
## 47	TheresePoletti	25	0.50
## 48	TimFarrand	39	0.78
## 49	ToddNissen	19	0.38
## 50	WilliamKazer	17	0.34

The model did well at predicting Jim Gilchrist - 100% accuracy. However, it did very poorly at predicting David Lawder. Let's take a look.

##	author	vector
## 1	AaronPressman	0
## 2	AlanCrosby	0
## 3	AlexanderSmith	0
## 4	BenjaminKangLim	0
## 5	BernardHickey	0
## 6	BradDorfman	8
## 7	DarrenSchuettler	0
## 8	DavidLawder	7
## 9	EdnaFernandes	0
## 10	EricAuchard	0
## 11	FumikoFujisaki	0
## 12	GrahamEarnshaw	0
## 13	HeatherScofield	0
## 14	JaneMacartney	0
## 15	JanLopatka	1
## 16	JimGilchrist	0
## 17	JoeOrtiz	0
## 18	JohnMastrini	1
## 19	JonathanBirt	0
## 20	JoWinterbottom	0
## 21	KarlPenhaul	1
## 22	KeithWeir	0
## 23	KevinDrawbaugh	3
## 24	KevinMorrison	0
## 25	KirstinRidley	0
## 26	KouroshKarimkhany	0

```
## 27      LydiaZajc      0
## 28      LynneO'Donnell 0
## 29      LynnleyBrowning 0
## 30      MarcelMichelson 0
## 31      MarkBendeich 0
## 32      MartinWolk     2
## 33      MatthewBunce 0
## 34      MichaelConnor 0
## 35      MureDickie     0
## 36      NickLouth      0
## 37      PatriciaCommings 0
## 38      PeterHumphrey 0
## 39      PierreTran     0
## 40      RobinSidel     4
## 41      RogerFillion 0
## 42      SamuelPerry    0
## 43      SarahDavison   0
## 44      ScottHillis    0
## 45      SimonCowell    0
## 46      TanEeLyn       0
## 47      TheresePoletti 0
## 48      TimFarrand     0
## 49      ToddNissen     23
## 50      WilliamKazer   0
```

The model attributed many of David Lawder's documents to the author Todd Nissen, so perhaps these two authors are difficult to distinguish.

Overall, the Naive Bayes model performed the best, and even though it assumes independent features and we know that many words are correlated with each other, the Naive Bayes model had much better accuracy than random forest so we will choose Naive Bayes.

Association Rules Mining

Here we will examine some interesting association rules among shopping baskets from the data on grocery purchases. We will set a support threshold of 0.01 to get rules that occurred in 1% of the data, and confidence of 0.5 to get rules that were correct at least half the time. Setting the maximum size to 4 made no difference so we set it to 3. This gave 15 rules.

```
inspect(groceryrules)
```

##	lhs	rhs	support	confidence	lift	count
## [1]	{curd,	=> {whole milk}	0.01006609	0.5823529	2.279125	99
##	yogurt}					
## [2]	{butter,	=> {whole milk}	0.01148958	0.5736041	2.244885	113
##	other vegetables}					
## [3]	{domestic eggs,	=> {whole milk}	0.01230300	0.5525114	2.162336	121
##	other vegetables}					
## [4]	{whipped/sour cream,	=> {whole milk}	0.01087951	0.5245098	2.052747	107
##	yogurt}					
## [5]	{other vegetables,	=> {whole milk}	0.01464159	0.5070423	1.984385	144
##	whipped/sour cream}					
## [6]	{other vegetables,	=> {whole milk}	0.01352313	0.5175097	2.025351	133
##	pip fruit}					


```
## [7] {citrus fruit,
##      root vegetables} => {other vegetables} 0.01037112 0.5862069 3.029608 102
## [8] {root vegetables,
##      tropical fruit}  => {other vegetables} 0.01230300 0.5845411 3.020999 121
## [9] {root vegetables,
##      tropical fruit}  => {whole milk}      0.01199797 0.5700483 2.230969 118
## [10] {tropical fruit,
##      yogurt}          => {whole milk}      0.01514997 0.5173611 2.024770 149
## [11] {root vegetables,
##      yogurt}          => {other vegetables} 0.01291307 0.5000000 2.584078 127
## [12] {root vegetables,
##      yogurt}          => {whole milk}      0.01453991 0.5629921 2.203354 143
## [13] {rolls/buns,
##      root vegetables} => {other vegetables} 0.01220132 0.5020921 2.594890 120
## [14] {rolls/buns,
##      root vegetables} => {whole milk}      0.01270971 0.5230126 2.046888 125
## [15] {other vegetables,
##      yogurt}          => {whole milk}      0.02226741 0.5128806 2.007235 219
```

Most of these rules predict whole milk, with a few predicting other vegetables. The rules predicting whole milk generally have other dairy items like yogurt, butter, or eggs, indicating people tend to buy milk when they buy other dairy items. Other vegetables were commonly associated with people buying root vegetables and some kind of fruit - these shoppers tend to buy fruits and vegetables together. Overall this gave a pretty small sample of rules, so we will change a couple parameters to try to look at other common rules.

We'll set the support to 0.001 to try and include more rules, and to compensate this we'll set confidence to 0.8 and only look at the rules with lift > 5. There was not much change increasing the maximum length above 5 so we set it to 5.

```
inspect(subset(groceryrules2, subset=lift > 5))
```

```
##      lhs                                rhs      support confidence      lift count
## [1] {liquor,
##      red/blush wine}                    => {bottled beer}    0.001931876 0.9047619 11.235269    19
## [2] {citrus fruit,
##      root vegetables,
##      soft cheese}                      => {other vegetables} 0.001016777 1.0000000 5.168156    10
## [3] {citrus fruit,
##      fruit/vegetable juice,
##      grapes}                          => {tropical fruit}   0.001118454 0.8461538 8.063879    11
## [4] {butter milk,
##      other vegetables,
##      pastry}                          => {yogurt}          0.001220132 0.8000000 5.734694    12
## [5] {pip fruit,
##      sausage,
##      sliced cheese}                    => {yogurt}          0.001220132 0.8571429 6.144315    12
## [6] {cream cheese,
##      margarine,
##      whipped/sour cream}              => {yogurt}          0.001016777 0.8333333 5.973639    10
## [7] {butter,
##      cream cheese,
##      root vegetables}                  => {yogurt}          0.001016777 0.9090909 6.516698    10
## [8] {butter,
##      tropical fruit,
##      white bread}                      => {yogurt}          0.001118454 0.8461538 6.065542    11
## [9] {beef,
```

##	butter,						
##	tropical fruit}	=> {yogurt}	0.001016777	0.8333333	5.973639	10	
## [10]	{fruit/vegetable juice,						
##	pork,						
##	tropical fruit}	=> {yogurt}	0.001016777	0.8333333	5.973639	10	
## [11]	{brown bread,						
##	pip fruit,						
##	whipped/sour cream}	=> {other vegetables}	0.001118454	1.0000000	5.168156	11	
## [12]	{butter,						
##	margarine,						
##	tropical fruit}	=> {yogurt}	0.001118454	0.8461538	6.065542	11	
## [13]	{fruit/vegetable juice,						
##	pastry,						
##	whipped/sour cream}	=> {yogurt}	0.001220132	0.8000000	5.734694	12	
## [14]	{other vegetables,						
##	rice,						
##	whole milk,						
##	yogurt}	=> {root vegetables}	0.001321810	0.8666667	7.951182	13	
## [15]	{grapes,						
##	tropical fruit,						
##	whole milk,						
##	yogurt}	=> {other vegetables}	0.001016777	1.0000000	5.168156	10	
## [16]	{ham,						
##	pip fruit,						
##	tropical fruit,						
##	yogurt}	=> {other vegetables}	0.001016777	1.0000000	5.168156	10	
## [17]	{ham,						
##	other vegetables,						
##	pip fruit,						
##	yogurt}	=> {tropical fruit}	0.001016777	0.8333333	7.941699	10	
## [18]	{ham,						
##	pip fruit,						
##	tropical fruit,						
##	whole milk}	=> {other vegetables}	0.001118454	1.0000000	5.168156	11	
## [19]	{butter,						
##	sliced cheese,						
##	tropical fruit,						
##	whole milk}	=> {yogurt}	0.001016777	0.9090909	6.516698	10	
## [20]	{oil,						
##	other vegetables,						
##	tropical fruit,						
##	whole milk}	=> {root vegetables}	0.001321810	0.8666667	7.951182	13	
## [21]	{cream cheese,						
##	curd,						
##	other vegetables,						
##	whipped/sour cream}	=> {yogurt}	0.001016777	0.9090909	6.516698	10	
## [22]	{cream cheese,						
##	curd,						
##	whipped/sour cream,						
##	whole milk}	=> {yogurt}	0.001118454	0.8461538	6.065542	11	
## [23]	{butter,						
##	other vegetables,						
##	tropical fruit,						
##	white bread}	=> {yogurt}	0.001016777	0.9090909	6.516698	10	

```

## [24] {beef,
##      citrus fruit,
##      other vegetables,
##      tropical fruit}      => {root vegetables}  0.001016777  0.8333333  7.645367    10
## [25] {butter,
##      curd,
##      other vegetables,
##      tropical fruit}      => {yogurt}          0.001016777  0.8333333  5.973639    10
## [26] {butter,
##      curd,
##      tropical fruit,
##      whole milk}          => {yogurt}          0.001220132  0.8571429  6.144315    12
## [27] {margarine,
##      root vegetables,
##      tropical fruit,
##      whole milk}          => {yogurt}          0.001016777  0.8333333  5.973639    10
## [28] {butter,
##      fruit/vegetable juice,
##      tropical fruit,
##      whipped/sour cream}  => {other vegetables} 0.001016777  1.0000000  5.168156    10
## [29] {newspapers,
##      rolls/buns,
##      soda,
##      whole milk}          => {other vegetables} 0.001016777  1.0000000  5.168156    10
## [30] {citrus fruit,
##      fruit/vegetable juice,
##      other vegetables,
##      soda}                => {root vegetables} 0.001016777  0.9090909  8.340400    10
## [31] {citrus fruit,
##      root vegetables,
##      tropical fruit,
##      whipped/sour cream}  => {other vegetables} 0.001220132  1.0000000  5.168156    12

```

This gives us some more interesting rules to work with. There are quite a few associations of yogurt here with other dairy items so that reaffirms the previous finding about dairy shoppers, as well as different fruits showing up in many of the rules that predict yogurt - makes sense as people like to mix fruit and yogurt. One very interesting rule here is liquor and red wine predicting bottled beer, as it has the highest lift and a very high confidence of over 90%. Basically it's very significant and people will likely buy bottled beer when they have bought liquor and wine as well, so a grocery store should be sure to put beer close by the liquor and wine, and maybe include some promotional marketing or coupons to maximize profits.